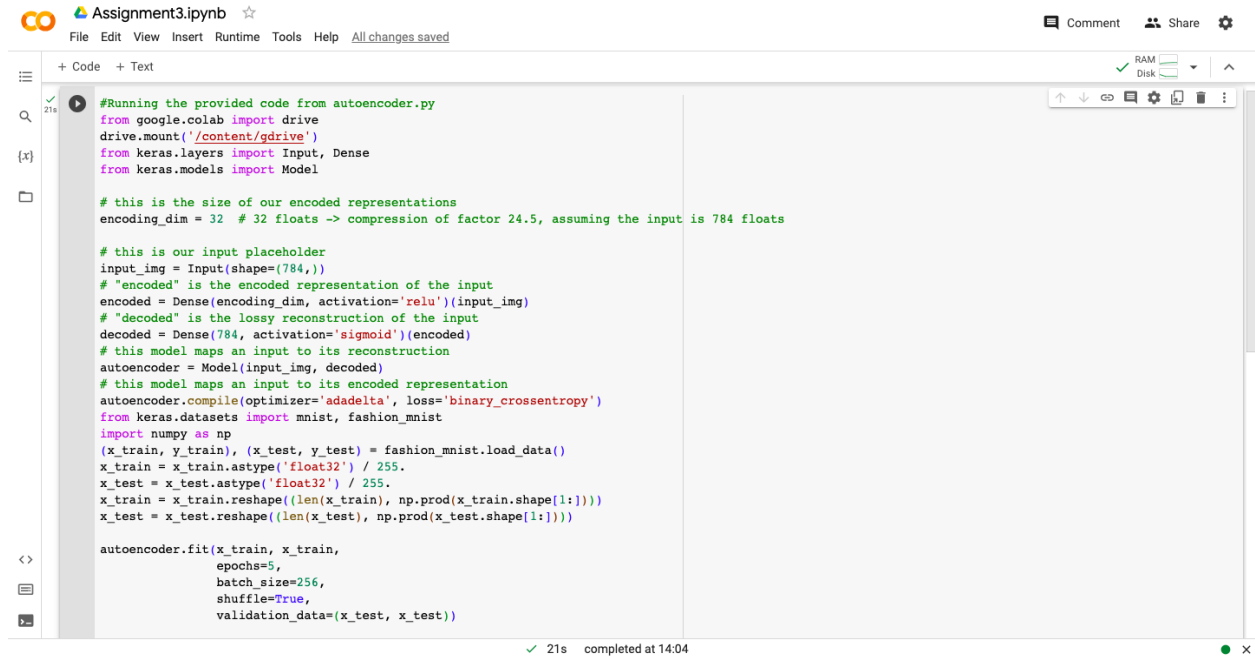


NN&DeepLearning_Lesson: Autoencoders

Running the provided autoencoder.py code



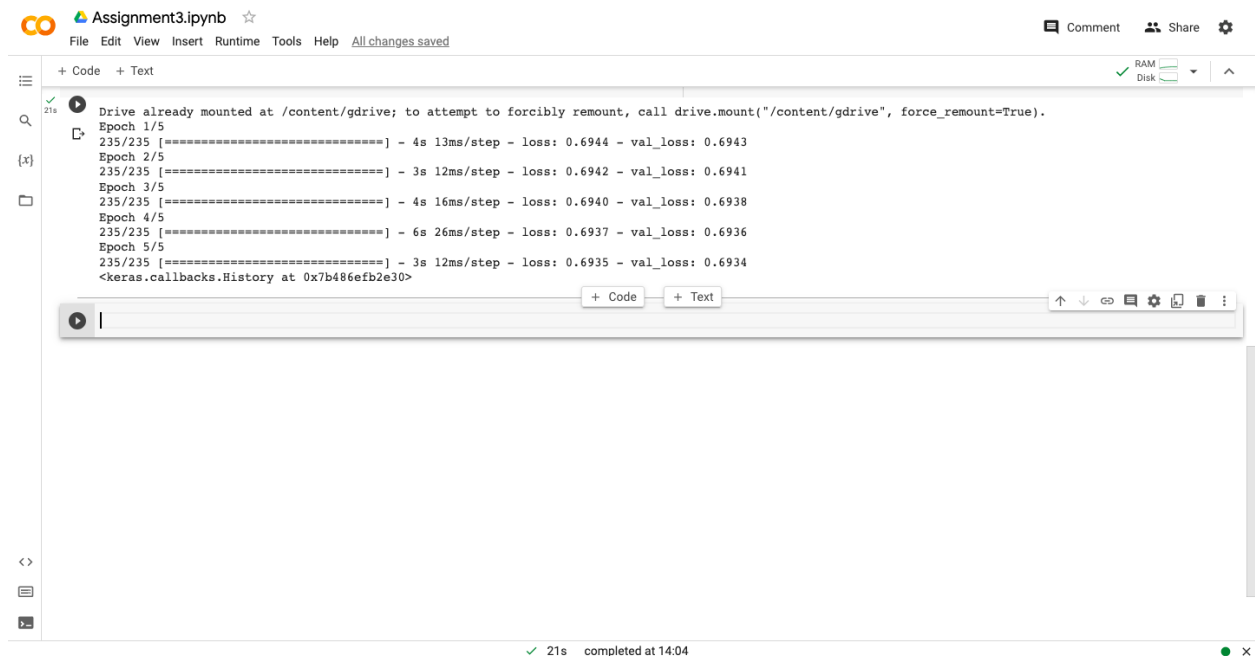
```
#Running the provided code from autoencoder.py
from google.colab import drive
drive.mount('/content/gdrive')
from keras.layers import Input, Dense
from keras.models import Model

# this is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# this is our input placeholder
input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
from keras.datasets import mnist, fashion_mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

autoencoder.fit(x_train, x_train,
                epochs=5,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test))
```

Output of the code is



```
Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).
Epoch 1/5
235/235 [=====] - 4s 13ms/step - loss: 0.6944 - val_loss: 0.6943
Epoch 2/5
235/235 [=====] - 3s 12ms/step - loss: 0.6942 - val_loss: 0.6941
Epoch 3/5
235/235 [=====] - 4s 16ms/step - loss: 0.6940 - val_loss: 0.6938
Epoch 4/5
235/235 [=====] - 6s 26ms/step - loss: 0.6937 - val_loss: 0.6936
Epoch 5/5
235/235 [=====] - 3s 12ms/step - loss: 0.6935 - val_loss: 0.6934
<keras.callbacks.History at 0x7b486efb2e30>
```

1. Add one more hidden layer to autoencoder

Ans : Code after adding the hidden layer is

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
#1.Added one more hidden layer to autoencoder

from keras.layers import Input, Dense
from keras.models import Model

# This is the size of our encoded representation
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# This is our input placeholder
input_img = Input(shape=(784,))

# "encoded" is the encoded representation of the input
encoded1 = Dense(128, activation='relu')(input_img)
encoded2 = Dense(encoding_dim, activation='relu')(encoded1)

# "decoded" is the lossy reconstruction of the input
decoded1 = Dense(128, activation='relu')(encoded2)
decoded2 = Dense(784, activation='sigmoid')(decoded1)

# This model maps an input to its reconstruction
autoencoder = Model(input_img, decoded2)

# This model maps an input to its encoded representation
encoder = Model(input_img, encoded2)

# This is our decoder model
encoded_input = Input(shape=(encoding_dim,))
decoder_layer1 = autoencoder.layers[-2]
decoder_layer2 = autoencoder.layers[-1]
decoder = Model(encoded_input, decoder_layer2(decoder_layer1(encoded_input)))

# Compile the model
autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy')

44s completed at 14:11
```

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
# Compile the model
autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy')

# Load the MNIST dataset
from keras.datasets import mnist, fashion_mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

# Normalize and flatten the data
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

# Train the autoencoder
autoencoder.fit(x_train, x_train,
               epochs=5,
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test))
```

Output is :

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Epoch 1/5
235/235 [=====] - 10s 36ms/step - loss: 0.6935 - val_loss: 0.6934
Epoch 2/5
235/235 [=====] - 7s 29ms/step - loss: 0.6933 - val_loss: 0.6933
Epoch 3/5
235/235 [=====] - 7s 31ms/step - loss: 0.6932 - val_loss: 0.6932
Epoch 4/5
235/235 [=====] - 8s 32ms/step - loss: 0.6931 - val_loss: 0.6930
Epoch 5/5
235/235 [=====] - 7s 29ms/step - loss: 0.6930 - val_loss: 0.6929
<keras.callbacks.History at 0x7b486f7f1fc0>
```

2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib.

Ans:

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

#2.Do the prediction on the test data and then visualize one of the reconstructed version of that test data.
#Also, visualize the same test data before reconstruction using Matplotlib

import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test)

# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))

# Plot the original image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Original Image")

# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()
```

Output is:



3. Repeat the question 2 on the denoising autoencoder

Running the code provided in Denoising_Autoencoder.py followed by the output :

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Reconnect ^

#Running the code provided in Denoising_Autoencoder.py
from keras.layers import Input, Dense
from keras.models import Model

# this is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# this is our input placeholder
input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics='accuracy')
from keras.datasets import fashion_mnist
import numpy as np
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

#introducing noise
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

autoencoder.fit(x_train_noisy, x_train,
               epochs=10,
               batch_size=256,
```

✓ 56s completed at 14:49

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Reconnect ^

autoencoder.fit(x_train_noisy, x_train,
               epochs=10,
               batch_size=256,
               shuffle=True,
               validation_data=(x_test_noisy, x_test_noisy))

Epoch 1/10
235/235 [=====] - 4s 16ms/step - loss: 0.6962 - accuracy: 0.0014 - val_loss: 0.6960 - val_accuracy: 0.0014
Epoch 2/10
235/235 [=====] - 4s 18ms/step - loss: 0.6960 - accuracy: 0.0014 - val_loss: 0.6957 - val_accuracy: 0.0014
Epoch 3/10
235/235 [=====] - 5s 22ms/step - loss: 0.6957 - accuracy: 0.0014 - val_loss: 0.6954 - val_accuracy: 0.0015
Epoch 4/10
235/235 [=====] - 4s 16ms/step - loss: 0.6954 - accuracy: 0.0014 - val_loss: 0.6952 - val_accuracy: 0.0017
Epoch 5/10
235/235 [=====] - 4s 15ms/step - loss: 0.6951 - accuracy: 0.0014 - val_loss: 0.6949 - val_accuracy: 0.0017
Epoch 6/10
235/235 [=====] - 5s 21ms/step - loss: 0.6949 - accuracy: 0.0014 - val_loss: 0.6947 - val_accuracy: 0.0018
Epoch 7/10
235/235 [=====] - 4s 19ms/step - loss: 0.6946 - accuracy: 0.0014 - val_loss: 0.6944 - val_accuracy: 0.0018
Epoch 8/10
235/235 [=====] - 4s 17ms/step - loss: 0.6944 - accuracy: 0.0014 - val_loss: 0.6942 - val_accuracy: 0.0019
Epoch 9/10
235/235 [=====] - 4s 18ms/step - loss: 0.6941 - accuracy: 0.0014 - val_loss: 0.6939 - val_accuracy: 0.0019
Epoch 10/10
235/235 [=====] - 8s 33ms/step - loss: 0.6939 - accuracy: 0.0014 - val_loss: 0.6937 - val_accuracy: 0.0019
<keras.callbacks.History at 0x7b486d0b640>
```

Ans: Now repeating all the steps from question 2 on the denoising autoencoder

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
#3.Repeat the question 2 on the denoising autoencoder

import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test_noisy)

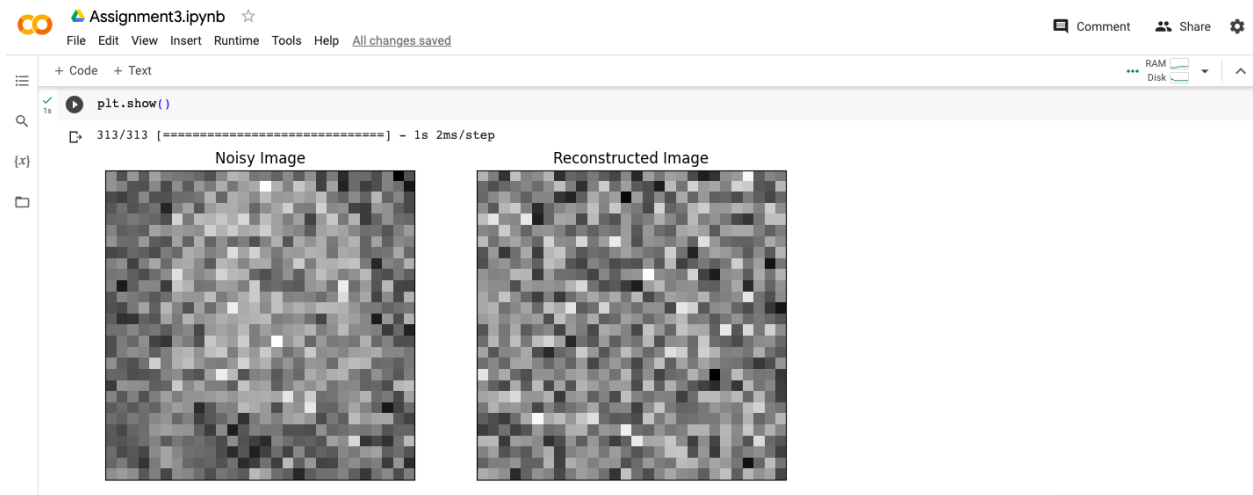
# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))

# Plot the original noisy image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test_noisy[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Noisy Image")

# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()
```

Output:



4. Plot loss and accuracy using the history object

Ans: Code is

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
#4.plot loss and accuracy using the history object

import matplotlib.pyplot as plt

# Train the autoencoder
history = autoencoder.fit(x_train_noisy, x_train,
                          epochs=10,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test_noisy, x_test_noisy))

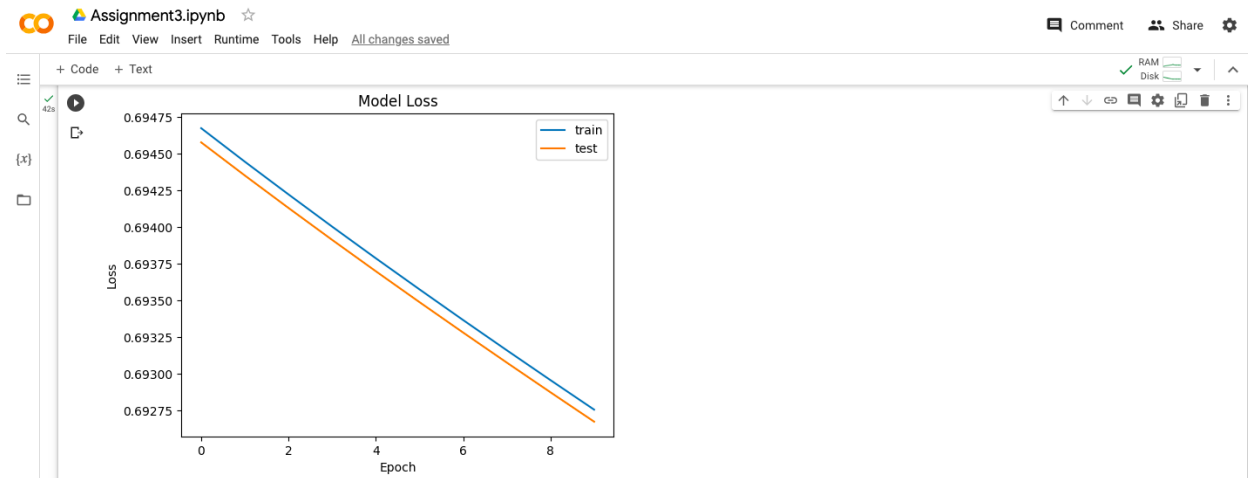
# Plot the loss
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()
plt.show()

# Plot the accuracy
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```

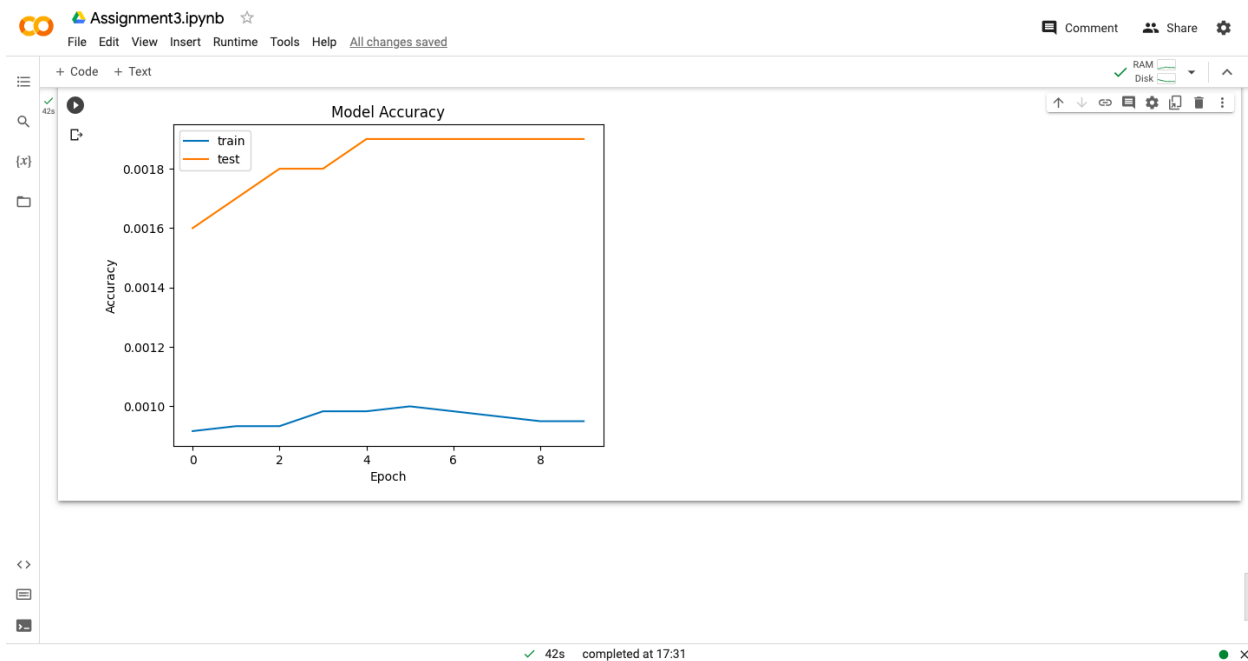
Output:

```
Assignment3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Epoch 1/10
235/235 [=====] - 3s 12ms/step - loss: 0.6947 - accuracy: 9.1667e-04 - val_loss: 0.6946 - val_accuracy: 0.0016
Epoch 2/10
235/235 [=====] - 3s 13ms/step - loss: 0.6944 - accuracy: 9.3333e-04 - val_loss: 0.6944 - val_accuracy: 0.0017
Epoch 3/10
235/235 [=====] - 3s 15ms/step - loss: 0.6942 - accuracy: 9.3333e-04 - val_loss: 0.6941 - val_accuracy: 0.0018
Epoch 4/10
235/235 [=====] - 3s 11ms/step - loss: 0.6940 - accuracy: 9.8333e-04 - val_loss: 0.6939 - val_accuracy: 0.0018
Epoch 5/10
235/235 [=====] - 3s 11ms/step - loss: 0.6938 - accuracy: 9.8333e-04 - val_loss: 0.6937 - val_accuracy: 0.0019
Epoch 6/10
235/235 [=====] - 3s 11ms/step - loss: 0.6936 - accuracy: 0.0010 - val_loss: 0.6935 - val_accuracy: 0.0019
Epoch 7/10
235/235 [=====] - 4s 16ms/step - loss: 0.6934 - accuracy: 9.8333e-04 - val_loss: 0.6933 - val_accuracy: 0.0019
Epoch 8/10
235/235 [=====] - 3s 12ms/step - loss: 0.6932 - accuracy: 9.6667e-04 - val_loss: 0.6931 - val_accuracy: 0.0019
Epoch 9/10
235/235 [=====] - 3s 11ms/step - loss: 0.6930 - accuracy: 9.5000e-04 - val_loss: 0.6929 - val_accuracy: 0.0019
Epoch 10/10
235/235 [=====] - 3s 11ms/step - loss: 0.6928 - accuracy: 9.5000e-04 - val_loss: 0.6927 - val_accuracy: 0.0019
```

Model Loss Plot:



Model Accuracy plot:



Github repo link:

https://github.com/LahariKollipara/NNDL_Assignment4